

The Role of Self-Reflection in Reinforcement Learning Agents

Dhruv Amin, Franco Bonilla, Arnav Chandra,
Dev Desai, Xinbei Miao, Myla Yambao, Nolan Zurek

Department of Computing Science

University of Alberta

Edmonton, Alberta, T6G 2E8, Canada

{dhruvs1, fgbonill, chandra3, devrajes, xinbei4, myambao, ntzurek}@ualberta.ca

July 11, 2026

Abstract

Self-reflection is the ability of an agent to consider its past actions or aspects of its own policy when making decisions; we prioritize the consideration of past actions. This paper investigates the role of self-reflection in reinforcement learning by modeling reflection as an explicit action for a tabular Q-learning agent. When the agent selects the reflect action, the environment state remains unchanged, a reflection cost is subtracted from the reward at that step, and the agent’s exploration and learning rates are updated based on recent performance. This approach incorporates the cost of reflection directly into the Markov Decision Process, facilitating controlled comparisons between a standard Q-learning baseline and reflective agents subjected to varying reflection costs. A 20×20 GridWorld navigation task is used to evaluate episodic return and reflection frequency, thereby identifying conditions under which reflection enhances learning efficiency and examining how reflection cost affects the agent’s willingness to reflect.

1 Introduction

Reinforcement learning (RL) studies how an agent can learn to behave optimally through trial and error by maximizing expected cumulative reward in an environment. In standard RL, the agent interacts with the environment over a sequence of states, selects actions, receives rewards, and updates its policy accordingly. This project investigates the role of self-reflection in a controlled RL setting. In our formulation, reflection is modeled as an explicit action available to the agent. When the agent executes the reflect action, the environment state does not change;

instead, the agent performs an internal update that modifies its learning parameters and incurs an immediate cost, which is subtracted from the reward. This penalty represents the cost of allocating a decision step to introspection rather than to environmental interaction, creating a direct trade-off between immediate and potential long-term benefit. Our goal is to determine whether introducing this reflection action leads to a net improvement in learning outcomes relative to a standard tabular Q-learning baseline. To isolate the impact of reflection, we compare two agents that are identical except for the availability and use of the reflect action. The baseline agent applies tabular Q-learning with an ϵ -greedy exploration strategy and fixed learning parameters, whereas the reflective agent may choose to reflect; upon reflecting, it adjusts internal parameters based on recent performance. We evaluate the reflective agent under varying reflection costs in a GridWorld environment, measuring episodic return and reflection frequency. These metrics directly address two key questions: under what conditions does self-reflection improve performance despite its cost, and how does the cost of reflection influence the agent’s decision to reflect. We hypothesize that in environments where standard directional actions are sufficient to learn an effective policy, reflection may impose an unnecessary penalty. In contrast, in settings where learning is unstable or requires adjustment, reflection may improve performance by guiding exploration and learning-rate adaptation.

2 Problem Formulation

We frame the agent’s interaction with the environment as a finite, episodic, Markov Decision Process (MDP)

$$M = (S, A, P, R, \gamma)$$

where S denotes the finite state space, A the set of available actions, $P(s' = s|s, a)$ the transition dynamics, $R(s, a, s')$ the reward function, and $\gamma \in [0, 1)$ the discount factor. The agent seeks to learn a policy $\pi(a|s)$ that maximizes the expected return over an episode. To incorporate self-reflection into this MDP, we include a reflection action, so the action set becomes

$$A = \{\text{up, down, left, right, reflect}\},$$

when the agent chooses the reflect action, the environment state remains unchanged $P(s' = s|s, a = \text{reflect}) = 1$. Reflection, however, incurs an immediate penalty that reduces the reward received on that timestep. This cost is implemented as $r \leftarrow r - c$ where $c \geq 0$ is a scalar reflection cost.

Let G_{reflect} denote the episodic return of a reflective agent and G_{baseline} the return of the baseline agent, without the ability to reflect. We want to determine under what conditions,

$$\mathbb{E}[G_{\text{reflect}}] > \mathbb{E}[G_{\text{baseline}}].$$

and how the reflection cost c influences this relationship. We hypothesize that when the transition and reward structure allows the agent to improve its return through selecting direction actions alone, reflection may impose an unnecessary cost. In contrast, environments that demand greater adaptation may benefit from reflection.

3 Related Work

In general, we found that definitions of reflection and self-reflection were not standardized. As such, much of the related work was consulted to learn about different formulations of reflection and how these formulations were evaluated. Much of this work seeks to solve problems that are not directly related to ours. In particular, a large part of the related work involved studying various notions of reflection in large language models.

Cox (2005) broadly surveys the field of metacognition and proposes various definitions for reflection. However, they do not formalize the cost of reflection

as it relates to selecting actions. Reflection is also assumed to always be beneficial to cognitive systems; this assumption is not tested empirically.

Omidshafiei et al. (2018) allow agents to learn when and what advice to give to peers to accelerate group learning. Under their formulation, a teacher agent is allowed to send advice signals to student agents in a multi-agent reinforcement-learning setting. The study shows that judicious teaching leads to a significant speedup in cooperative learning. However, the formulation treats advice as an arbitrary signal rather than a self-articulation of an agent’s own policy. A cost for giving or receiving advice is not modelled. The formulation also assumes advice is generally beneficial and does not analyze the possible performance decrease from receiving badly formulated or misguided advice.

Renze and Guven (2024) examine the impact of learning from past traces on the problem-solving performance of large language model agents. By comparing baseline prompts with reflection-enabled prompts across a range of reasoning tasks, they show that reflection often increases accuracy and helps detect mistakes in single-episode reasoning. However, this work is limited to online or single-agent language tasks and does not address sequential decision making with explicit episodic returns. In addition, they do not explicitly model a cost for producing reflections such as computational, time, or energy penalty. Finally, they do not explore how self-reflection may influence other agents or propagate errors in a multi-agent setting.

Zhang et al. (2024) propose using reflection as a mechanism to adjust the internal policy components of LLMs to improve performance after failure. Under their formulation, an LLM agent articulates beliefs about the world and itself in natural language based on external performance feedback and its previous beliefs. These are included as part of the subsequent prompt. In this way, the articulations act like a natural language extension of the agent’s policy. Thus, the agent considers aspects of its own policy as part of reflection. However, their model does not assign a cost to reflection, nor does it explore when reflection may degrade the performance.

Taken together, these lines of research show that various forms of self-reflection or advice can improve performance. However, most research in this area does not address the cost-benefit trade-offs and fidelity of choosing to self-reflect.

4 Proposed Approach

Our aim is to determine whether adding an explicit reflection action to a standard reinforcement-learning agent improves learning efficiency despite its immediate cost. Following the problem formulation in Section 2, we structure the experiments to isolate reflection cost as the only differing component between agents, and use a clearly defined baseline for comparison. We extend a tabular Q-learning agent with a reflect action, that leaves the environment state unchanged, but modifies the agent’s internal learning parameters while incurring a fixed penalty. This formulation allows us to evaluate the effect of reflection on learning performance, and if the benefit produced by reflection outweighs its immediate cost.

The objective is to determine whether incorporating an explicit reflection action into a standard reinforcement-learning agent enhances learning efficiency despite its immediate cost. Consistent with the problem formulation in Section 2, the experiments are structured to ensure that reflection cost is the sole variable distinguishing agents, with a clearly defined baseline for comparison. A tabular Q-learning agent is extended with a reflect action, which leaves the environment state unchanged but modifies the agent’s internal learning parameters while incurring a fixed penalty. This approach facilitates assessment of the impact of reflection on learning performance and evaluation of whether the benefits of reflection outweigh its immediate costs.

4.1 Baseline Agent

The baseline agent is a standard tabular Q-learning agent with an ϵ -greedy exploration strategy. It receives no other guidance or heuristics. The baseline agent updates its action-value estimates using the Q-learning rule

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)].$$

where α is the learning rate initialized as 0.3, γ is the discount factor initialized as 0.95, and r is the reward obtained after executing action a in state s .

4.2 Reflective Agent

The reflective agent is identical to the baseline in its use of tabular Q-learning, with the exception that it may choose an additional reflect action. Executing the reflect action leaves the environment unchanged

but imposed an immediate penalty on the reward implemented as

$$r \leftarrow r - c$$

where $c \geq 0$ and denotes the reflection cost defined in the MDP. We evaluate three cost levels, where $c = \{0.1, 0.2, 0.3\}$, to examine whether the cost of the reflection action has an influence on the agent’s decision to reflect. Upon reflecting, the agent updates its internal learning parameters, specifically the exploration rate ϵ and the learning rate α . Specifically, the agent compares the average reward over its most recent five actions with the average reward from the preceding five. If the performance has improved, ϵ is decreased by 0.02 to a value no smaller than 0.01 and α is increased by 0.02 to a value no larger than 1, therefore promoting more exploitative behavior and placing greater weight on the new information. If performance has deteriorated, ϵ is increased and α is decreased, allowing for more exploration while reducing the influence of potentially misleading rewards.

4.3 GridWorld Environment

The experiments are conducted in a 20×20 GridWorld environment in which start and goal positions, as well as obstacle placements, are randomly placed to produce different maps. To generate variation in navigational complexity, we construct maps with three structural configurations: 100 obstacles with a minimum start-goal distance of 10, 150 obstacles with a minimum distance of 15, and 200 obstacles with a minimum distance of 20. We classify the maps into different degrees of difficulty based on baseline agent performance. For each map configuration, the baseline agent is trained for three independent runs of 1000 episodes, and the resulting success rates are averaged to obtain a performance measure x . Maps for which the baseline achieves an average success rate greater than 95%, $x \geq 95\%$ are labeled easy, those with $25\% \leq x < 95\%$ are labeled medium, and those with $x < 25\%$ are designated hard.




Figure 1: An example instance of an environment. Red squares are obstacles. The goal state is green and the start state is blue.

4.4 Experiment

We train three reflective agents, each with differing reflection costs, and the baseline agent on identical sequences of GridWorld environments. All agents interact with the same transition dynamics: movement actions advance the agent deterministically unless blocked by an obstacle, each step incurs a reward of -0.1 , and reaching the goal yields a terminal reward of $+10$. Each agent is evaluated over 300 episodes, with a maximum of 10000 steps per episode. Throughout training, we record episodic return achieved by each agent as well as the number of times the agent chooses the reflect action. These metrics are collected across all map difficulty levels.

5 Empirical Evaluation

To evaluate the utility of self-reflection in reinforcement learning, we check whether reflection improves the performance of Q-learning on grid-based pathfinding environments. Specifically, we examine if the benefit of exploration and exploitation rate parameter adjustment outweighs the cost of reflection in returns and missed opportunities of movement. In this section, we present our results in terms of re-




Figure 2: Average Return (Easy)

flection frequency, α and ϵ over time, episodic returns and success rate for each agent. The goal of this evaluation is to find which types of environments self-reflection leads to the greatest improvement in performance, where it is counterproductive, and how this trade-off changes across environments of varying complexity.

5.1 Performance by Difficulty Level

5.1.1 Easy Maps

In easy environments, reflection provided no measurable benefit. The low-cost reflective agent had the most flections, while the medium-cost agent reflected moderately and the high-cost agent rarely chose to reflect. The baseline and high-cost agents obtained the highest success rates and the most favorable episodic returns, with the medium-cost agent performing similarly but not better than the baseline. The low-cost agent, despite frequent early reflection, achieved the lowest episodic returns, indicating that the return cost of reflection and missed movement opportunity outweighed any learning advantage. All reflective agents learned over time to reduce the frequency of reflection, but this reduction did not translate into improvements over the baseline. These results suggest that in environments that are already easily solvable, reflection introduces unnecessary cost and does not provide strategic insight that improves learning outcomes.

5.1.2 Medium Maps

In medium environments, the benefits of reflection started to show up, but behaved differently across




Figure 3: Average Reflect (Easy)




Figure 4: Average Epsilon (Easy)




Figure 5: Average Alpha (Easy)




Figure 6: Average Return (Medium)




Figure 7: Average Reflect (Medium)

cost levels. The low-cost reflective agent stopped reflecting earlier than the medium- and high-cost agents, while the high-cost agent continued reflecting throughout training, likely due to difficulty recovering from the more extreme adjustments to exploration and exploitation. Both low- and medium-cost agents achieved successful episodes, in contrast to the baseline agent, which did not solve the task within the allotted steps significantly. These results suggest that reflection can help to overcome challenges with minimal rewards, although excessive cost may prevent the agent's ability to escape suboptimal behavior once reflection is repeatedly punished.

5.1.3 Hard Maps

Reflection demonstrated its greatest utility in high-difficulty environments. Both the low- and medium-cost agents achieved the highest success rates and




Figure 8: Average Epsilon (Medium)




Figure 9: Average Alpha (Medium)




Figure 10: Average Return (Hard)




Figure 11: Average Reflect (Hard)

episodic returns, substantially outperforming the baseline agent, which failed to complete the task in the given training horizon on average. The low-cost agent reflected frequently in the early episodes but gradually reduced reflection once a reliable path to the goal was discovered, leading to consistently higher returns. The medium-cost agent followed a similar pattern, though with fewer total reflections. These findings indicate that reflection can speed up learning in environments that are not trivial but rather complex.

5.1.4 Summary

Across map difficulties, reflection demonstrates value only when the environment is sufficiently complex and requires more deliberate exploration. In easy environments, reflection offers no benefit; agents quickly learn that revising decisions is unneces-




Figure 12: Average Epsilon (Hard)




Figure 13: Average Alpha (Hard)

sary, and performance decreases when reflection incurs avoidable cost. In medium and hard environments, reflection directly improves learning by enabling agents to increase or decrease their exploration and learning rate in order to find a solution path, leading to higher rewards and successful paths that the baseline almost never discovers. Reflection helps learning by speeding up correction of bad actions, reducing repeated mistakes, and enabling earlier convergence to useful behaviors. Adjusting α and ϵ speeds up convergence rate to a solution by tuning the learning rate and exploration rate to encourage such solutions.

6 Discussion

We have established a functioning single-agent environment in Gridworld where reflection is an additional action incurring a cost.

From the initial results, we can say that the environment and algorithm behave as expected. The agents learn to maximize reward while avoiding costly steps. However, because the reflection operator $F()$ is not yet performing actual self-analysis, the results currently capture only the agent’s initial learning behaviour before self-reflection is actually implemented rather than the intended cognitive trade-off.

Based on the results from the first version of our program, we can see that adding reflection as an explicit action is compatible with standard reinforcement-learning formulations and provides a flexible framework for extending to meta-learning experiments. However, these are only our early findings and interpretability is limited because reflection currently has no positive utility.

7 Shortcomings & Future Work

7.1 Shortcomings

The reflection mechanism is narrowly defined: reflection only increases the exploration rate after poor performance, a change that neither reliably improves learning nor meaningfully summarizes past behavior. Under this formulation, the agent has little reason to reflect, since the only modeled downside is the explicit reflection cost. Without any additional risks or long-term advantages tied to reflection, the agent cannot learn a compelling strategic use for the reflect action. Furthermore, the use of tabular Q-learning and coordinate-based state representations

limits both task complexity and the agent’s capacity to generalize. Because states encode only positions rather than structural features of the environment, the agent cannot connect successful reflection in one region of the map to similar situations elsewhere. Reflection, therefore, becomes beneficial only after the optimal path is already discovered—at which point the agent no longer needs to reflect, making the mechanism effectively redundant in practice.

7.2 Future Work

A key area of exploration would be to formulate and adopt a more robust model of reflection. In particular, we would allow agents to articulate and interpret aspects of their own policy in addition to their past actions. This formulation would encompass more of the widely-used definitions of reflection, leading to more general results.

This project could also be expanded to include multiple agents instead of just one. Each agent will have its own learning process and reflection steps, but they will share the same environment. Agents can either cooperate or complement each other while learning. This setting also raises the new possibility of shared reflection between agents by empowering agents to consider the past trajectories of other agents. By expanding our experiments to multi-agent settings, we would be able to better understand how reflection spreads between agents and how more opportunities for reflection affects overall learning speed.

Another goal for future work is allow agents to decide when to reflect and when to act. Agents would be able to choose whether to share their reflections or keep them private. This formulation could serve as a better model of the real-world behavior of robots or artificial intelligence systems.

8 Conclusions

This project investigated whether adding an explicit reflect action to a standard reinforcement-learning agent can improve performance despite an immediate cost. We compared a baseline tabular Q-learning agent that can only take directional actions (up, down, left, right) to a reflective variant that can additionally choose reflect. When the agent reflects, the environment state remains unchanged, the agent receives an immediate penalty c and the agent adjusts its internal learning behavior by updating the

exploration rate ϵ and learning rate α using a recent-performance comparison.

Across experiments in GridWorld, we evaluated episodic return and how frequently the agent chose to reflect under multiple reflection-cost settings. The results indicate that reflection is not universally beneficial: its utility depends on whether the adaptive updates to ϵ and α produce improvements large enough to outweigh the reflection cost. As the reflection cost increases, the agent has less incentive to choose the reflect action, so reflection becomes less frequent and any benefits must be large enough to offset the added penalty; when the cost is low, the penalty is smaller, making reflection more attractive and potentially beneficial if the α and/or ϵ adjustments improve subsequent decisions.

Two validation points are important for interpreting the findings. First, if the maps are categorized as easy, medium, and hard, these categories are verified empirically for the baseline agent to ensure they match actual difficulty. Second, to confirm that performance changes are truly caused by reflection’s internal parameter updates, an ablation where reflect incurs the same cost but does not modify α or ϵ would provide a clean control. Overall, this work supports the view that self-reflection is best understood as a cost-sensitive mechanism: it can be beneficial in regimes where adaptive learning is needed, but it can become neutral or detrimental when its cost outweighs its benefits.

Acknowledgments

We used AI assistance (specifically ChatGPT) to help check grammar and improve sentence structure in this document.

References

- Cox, M. T. (2005). Metacognition in computation: A selected history. In *Proceedings of the AAAI Spring Symposium on Metacognition in Computation*, Cambridge, MA. AAAI Press. BBN Technologies, 10 Moulton St., Cambridge, MA 02138.
- Omidshafiei, S., Kim, D.-K., Liu, M., Tesauro, G., Riemer, M., Amato, C., Campbell, M., and How, J. P. (2018). Learning to teach in cooperative multiagent reinforcement learning. *arXiv preprint arXiv:1805.07830*. Version 4, [cs.MA], 31 Aug 2018.

Renze, M. and Guven, E. (2024). Self-reflection in llm agents: Effects on problem-solving performance. *arXiv preprint arXiv:2405.06682*. Version 3, [cs.CL], 16 Oct 2024.

Zhang, W., Tang, K., Wu, H., Wang, M., Shen, Y., Hou, G., Tan, Z., Li, P., Zhuang, Y., and Lu, W. (2024). Agent-pro: Learning to evolve via policy-level reflection and optimization. *arXiv preprint arXiv:*. Affiliations: Zhejiang University, Chinese Academy of Sciences, Nanjing University of Posts and Telecommunications, Beijing University of Technology, University of Chinese Academy of Sciences.